

Customer churn Prediction using ML

June 2, 2025

```
[3]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import LabelEncoder
from imblearn.over_sampling import SMOTE
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, \
    classification_report
import pickle
```

```
[4]: df = pd.read_csv('WA_Fn-UseC_-Telco-Customer-Churn.csv')
```

```
[6]: df.shape
```

```
[6]: (7043, 21)
```

```
[7]: df.head()
```

```
[7]:  customerID  gender  SeniorCitizen  Partner  Dependents  tenure  PhoneService  \
0  7590-VHVEG  Female              0      Yes           No         1           No
1  5575-GNVDE   Male              0      No            No        34           Yes
2  3668-QPYBK   Male              0      No            No         2           Yes
3  7795-CFOCW   Male              0      No            No        45           No
4  9237-HQITU   Female            0      No            No         2           Yes
```

```
    MultipleLines  InternetService  OnlineSecurity  ...  DeviceProtection  \
0  No phone service              DSL              No  ...              No
1                No              DSL              Yes  ...              Yes
2                No              DSL              Yes  ...              No
3  No phone service              DSL              Yes  ...              Yes
4                No  Fiber optic              No  ...              No
```

```
    TechSupport  StreamingTV  StreamingMovies  ...  Contract  PaperlessBilling  \
0            No            No              No  ...  Month-to-month              Yes
```

1	No	No	No	One year	No
2	No	No	No	Month-to-month	Yes
3	Yes	No	No	One year	No
4	No	No	No	Month-to-month	Yes

	PaymentMethod	MonthlyCharges	TotalCharges	Churn
0	Electronic check	29.85	29.85	No
1	Mailed check	56.95	1889.5	No
2	Mailed check	53.85	108.15	Yes
3	Bank transfer (automatic)	42.30	1840.75	No
4	Electronic check	70.70	151.65	Yes

[5 rows x 21 columns]

```
[8]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   customerID            7043 non-null   object
1   gender                 7043 non-null   object
2   SeniorCitizen          7043 non-null   int64
3   Partner                7043 non-null   object
4   Dependents             7043 non-null   object
5   tenure                 7043 non-null   int64
6   PhoneService           7043 non-null   object
7   MultipleLines          7043 non-null   object
8   InternetService        7043 non-null   object
9   OnlineSecurity         7043 non-null   object
10  OnlineBackup           7043 non-null   object
11  DeviceProtection       7043 non-null   object
12  TechSupport            7043 non-null   object
13  StreamingTV            7043 non-null   object
14  StreamingMovies        7043 non-null   object
15  Contract               7043 non-null   object
16  PaperlessBilling       7043 non-null   object
17  PaymentMethod          7043 non-null   object
18  MonthlyCharges         7043 non-null   float64
19  TotalCharges           7043 non-null   object
20  Churn                  7043 non-null   object
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB
```

```
[9]: # dropping customerID column as this is not required for modelling
df = df.drop(columns=["customerID"])
```

```
[10]: df.head(2)
```

```
[10]:   gender  SeniorCitizen  Partner  Dependents  tenure  PhoneService  \
0  Female              0      Yes          No         1           No
1   Male              0      No          No        34           Yes

      MultipleLines  InternetService  OnlineSecurity  OnlineBackup  \
0  No phone service              DSL              No           Yes
1              No              DSL              Yes           No

      DeviceProtection  TechSupport  StreamingTV  StreamingMovies  Contract  \
0              No              No              No              No  Month-to-month
1              Yes              No              No              No    One year

      PaperlessBilling  PaymentMethod  MonthlyCharges  TotalCharges  Churn
0              Yes  Electronic check          29.85          29.85    No
1              No    Mailed check          56.95         1889.5    No
```

```
[11]: df.columns
```

```
[11]: Index(['gender', 'SeniorCitizen', 'Partner', 'Dependents', 'tenure',
        'PhoneService', 'MultipleLines', 'InternetService', 'OnlineSecurity',
        'OnlineBackup', 'DeviceProtection', 'TechSupport', 'StreamingTV',
        'StreamingMovies', 'Contract', 'PaperlessBilling', 'PaymentMethod',
        'MonthlyCharges', 'TotalCharges', 'Churn'],
        dtype='object')
```

```
[12]: print(df["gender"].unique())
```

```
['Female' 'Male']
```

```
[13]: print(df["SeniorCitizen"].unique())
```

```
[0 1]
```

```
[14]: # printing the unique values in all the columns
```

```
numerical_features_list = ["tenure", "MonthlyCharges", "TotalCharges"]
```

```
for col in df.columns:
    if col not in numerical_features_list:
        print(col, df[col].unique())
        print("-"*50)
```

```
gender ['Female' 'Male']
```

```
-----
SeniorCitizen [0 1]
```

```
-----
Partner ['Yes' 'No']
```

```

-----
Dependents ['No' 'Yes']
-----
PhoneService ['No' 'Yes']
-----
MultipleLines ['No phone service' 'No' 'Yes']
-----
InternetService ['DSL' 'Fiber optic' 'No']
-----
OnlineSecurity ['No' 'Yes' 'No internet service']
-----
OnlineBackup ['Yes' 'No' 'No internet service']
-----
DeviceProtection ['No' 'Yes' 'No internet service']
-----
TechSupport ['No' 'Yes' 'No internet service']
-----
StreamingTV ['No' 'Yes' 'No internet service']
-----
StreamingMovies ['No' 'Yes' 'No internet service']
-----
Contract ['Month-to-month' 'One year' 'Two year']
-----
PaperlessBilling ['Yes' 'No']
-----
PaymentMethod ['Electronic check' 'Mailed check' 'Bank transfer (automatic)'
'Credit card (automatic)']
-----
Churn ['No' 'Yes']
-----

```

```
[15]: print(df.isnull().sum())
```

```

gender          0
SeniorCitizen   0
Partner         0
Dependents      0
tenure          0
PhoneService    0
MultipleLines   0
InternetService 0
OnlineSecurity  0
OnlineBackup    0
DeviceProtection 0
TechSupport     0
StreamingTV     0
StreamingMovies 0
Contract       0

```

```
PaperlessBilling    0
PaymentMethod       0
MonthlyCharges      0
TotalCharges        0
Churn               0
dtype: int64
```

```
[16]: #df["TotalCharges"] = df["TotalCharges"].astype(float)
df[df["TotalCharges"]==" "]
```

```
[16]:      gender  SeniorCitizen  Partner  Dependents  tenure  PhoneService  \
488   Female                0     Yes         Yes       0           No
753    Male                0     No          Yes       0           Yes
936   Female                0     Yes         Yes       0           Yes
1082   Male                0     Yes         Yes       0           Yes
1340  Female                0     Yes         Yes       0           No
3331   Male                0     Yes         Yes       0           Yes
3826   Male                0     Yes         Yes       0           Yes
4380  Female                0     Yes         Yes       0           Yes
5218   Male                0     Yes         Yes       0           Yes
6670  Female                0     Yes         Yes       0           Yes
6754   Male                0     No          Yes       0           Yes
```

```
      MultipleLines  InternetService      OnlineSecurity  \
488   No phone service          DSL          Yes
753                No          No  No internet service
936                No          DSL          Yes
1082                Yes          No  No internet service
1340  No phone service          DSL          Yes
3331                No          No  No internet service
3826                Yes          No  No internet service
4380                No          No  No internet service
5218                No          No  No internet service
6670                Yes          DSL          No
6754                Yes          DSL          Yes
```

```
      OnlineBackup  DeviceProtection      TechSupport  \
488                No          Yes          Yes
753  No internet service  No internet service  No internet service
936                Yes          Yes          No
1082  No internet service  No internet service  No internet service
1340                Yes          Yes          Yes
3331  No internet service  No internet service  No internet service
3826  No internet service  No internet service  No internet service
4380  No internet service  No internet service  No internet service
5218  No internet service  No internet service  No internet service
6670                Yes          Yes          Yes
```

	Yes	No	Yes
	StreamingTV	StreamingMovies	Contract PaperlessBilling \
488	Yes	No	Two year Yes
753	No internet service	No internet service	Two year No
936	Yes	Yes	Two year No
1082	No internet service	No internet service	Two year No
1340	Yes	No	Two year No
3331	No internet service	No internet service	Two year No
3826	No internet service	No internet service	Two year No
4380	No internet service	No internet service	Two year No
5218	No internet service	No internet service	One year Yes
6670	Yes	No	Two year No
6754	No	No	Two year Yes

	PaymentMethod	MonthlyCharges	TotalCharges	Churn
488	Bank transfer (automatic)	52.55		No
753	Mailed check	20.25		No
936	Mailed check	80.85		No
1082	Mailed check	25.75		No
1340	Credit card (automatic)	56.05		No
3331	Mailed check	19.85		No
3826	Mailed check	25.35		No
4380	Mailed check	20.00		No
5218	Mailed check	19.70		No
6670	Mailed check	73.35		No
6754	Bank transfer (automatic)	61.90		No

```
[17]: len(df[df["TotalCharges"]==" "])
```

```
[17]: 11
```

```
[19]: df["TotalCharges"] = df["TotalCharges"].replace({" ": "0.0"})
```

```
[20]: df["TotalCharges"] = df["TotalCharges"].astype(float)
```

```
[21]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 20 columns):
#   Column          Non-Null Count  Dtype
---  -
0   gender          7043 non-null   object
1   SeniorCitizen    7043 non-null   int64
2   Partner          7043 non-null   object
3   Dependents       7043 non-null   object
4   tenure          7043 non-null   int64
```

```

5   PhoneService      7043 non-null   object
6   MultipleLines     7043 non-null   object
7   InternetService   7043 non-null   object
8   OnlineSecurity    7043 non-null   object
9   OnlineBackup      7043 non-null   object
10  DeviceProtection  7043 non-null   object
11  TechSupport       7043 non-null   object
12  StreamingTV       7043 non-null   object
13  StreamingMovies   7043 non-null   object
14  Contract          7043 non-null   object
15  PaperlessBilling  7043 non-null   object
16  PaymentMethod     7043 non-null   object
17  MonthlyCharges    7043 non-null   float64
18  TotalCharges      7043 non-null   float64
19  Churn             7043 non-null   object
dtypes: float64(2), int64(2), object(16)
memory usage: 1.1+ MB

```

```
[22]: # checking the class distribution of target column
print(df["Churn"].value_counts())
```

```

Churn
No      5174
Yes     1869
Name: count, dtype: int64

```

1 Insights:

1)Customer Id removed as it is not required for modelling. # 2)No missing values in the dataset.
 # 3)Missing values in the TotalCharges column were replaced with 0. # 4)Class imbalance identified in the target .

2 Exploratory Data Analysis (EDA)

```
[23]: df.shape
```

```
[23]: (7043, 20)
```

```
[24]: df.columns
```

```

[24]: Index(['gender', 'SeniorCitizen', 'Partner', 'Dependents', 'tenure',
          'PhoneService', 'MultipleLines', 'InternetService', 'OnlineSecurity',
          'OnlineBackup', 'DeviceProtection', 'TechSupport', 'StreamingTV',
          'StreamingMovies', 'Contract', 'PaperlessBilling', 'PaymentMethod',
          'MonthlyCharges', 'TotalCharges', 'Churn'],
          dtype='object')

```

```
[25]: df.describe()
```

```
[25]:
```

	SeniorCitizen	tenure	MonthlyCharges	TotalCharges
count	7043.000000	7043.000000	7043.000000	7043.000000
mean	0.162147	32.371149	64.761692	2279.734304
std	0.368612	24.559481	30.090047	2266.794470
min	0.000000	0.000000	18.250000	0.000000
25%	0.000000	9.000000	35.500000	398.550000
50%	0.000000	29.000000	70.350000	1394.550000
75%	0.000000	55.000000	89.850000	3786.600000
max	1.000000	72.000000	118.750000	8684.800000

3 Numerical Feature - Analysis

```
[26]: def plot_histogram(df, column_name):

    plt.figure(figsize=(5, 3))
    sns.histplot(df[column_name], kde=True)
    plt.title(f"Distribution of {column_name}")


    # calculate the mean and median values for the columns
    col_mean = df[column_name].mean()
    col_median = df[column_name].median()

    # add vertical lines for mean and median
    plt.axvline(col_mean, color="red", linestyle="--", label="Mean")
    plt.axvline(col_median, color="green", linestyle="-", label="Median")


    plt.legend()

    plt.show()
```


```
[27]: plot_histogram(df, "tenure")
```

```
[28]: plot_histogram(df, "MonthlyCharges")
```



```
[29]: plot_histogram(df, "TotalCharges")
```




4 Box plot for numerical feature


```
[30]: def plot_boxplot(df, column_name):

    plt.figure(figsize=(5, 3))
    sns.boxplot(y=df[column_name])
    plt.title(f"Box Plot of {column_name}")
    plt.ylabel(column_name)
    plt.show
```

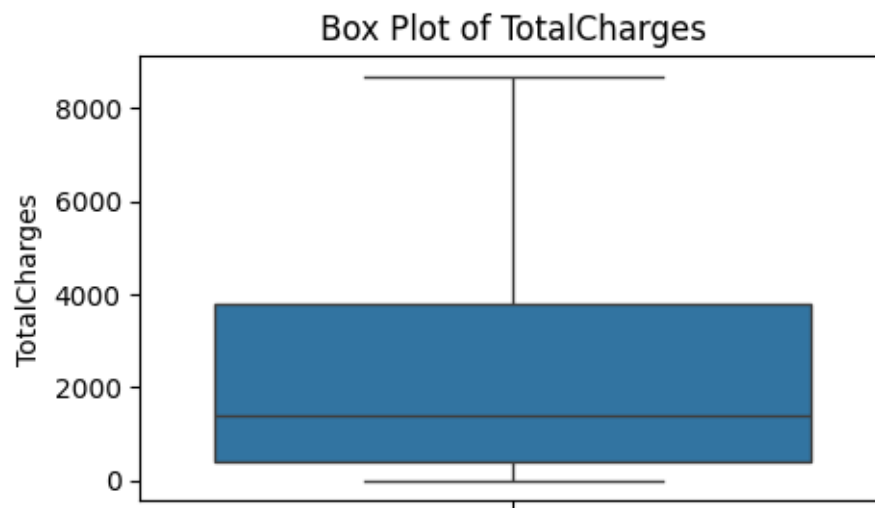
```
[31]: plot_boxplot(df, "tenure")
```



```
[33]: plot_boxplot(df, "MonthlyCharges")
```




```
[34]: plot_boxplot(df, "TotalCharges")
```



5 Correlation Heatmap for numerical columns

```
[35]: # correlation matrix - heatmap
plt.figure(figsize=(8, 4))
sns.heatmap(df[["tenure", "MonthlyCharges", "TotalCharges"]].corr(),
            annot=True, cmap="coolwarm", fmt=".2f")
plt.title("Correlation Heatmap")
plt.show()
```



6 Categorical feature - Analysis

```
[37]: df.columns
```

```
[37]: Index(['gender', 'SeniorCitizen', 'Partner', 'Dependents', 'tenure',
        'PhoneService', 'MultipleLines', 'InternetService', 'OnlineSecurity',
        'OnlineBackup', 'DeviceProtection', 'TechSupport', 'StreamingTV',
        'StreamingMovies', 'Contract', 'PaperlessBilling', 'PaymentMethod',
        'MonthlyCharges', 'TotalCharges', 'Churn'],
        dtype='object')
```

```
[38]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 20 columns):
```

#	Column	Non-Null Count	Dtype
0	gender	7043 non-null	object
1	SeniorCitizen	7043 non-null	int64
2	Partner	7043 non-null	object
3	Dependents	7043 non-null	object
4	tenure	7043 non-null	int64
5	PhoneService	7043 non-null	object
6	MultipleLines	7043 non-null	object
7	InternetService	7043 non-null	object
8	OnlineSecurity	7043 non-null	object
9	OnlineBackup	7043 non-null	object
10	DeviceProtection	7043 non-null	object
11	TechSupport	7043 non-null	object
12	StreamingTV	7043 non-null	object
13	StreamingMovies	7043 non-null	object
14	Contract	7043 non-null	object
15	PaperlessBilling	7043 non-null	object
16	PaymentMethod	7043 non-null	object
17	MonthlyCharges	7043 non-null	float64
18	TotalCharges	7043 non-null	float64
19	Churn	7043 non-null	object

dtypes: float64(2), int64(2), object(16)



memory usage: 1.1+ MB

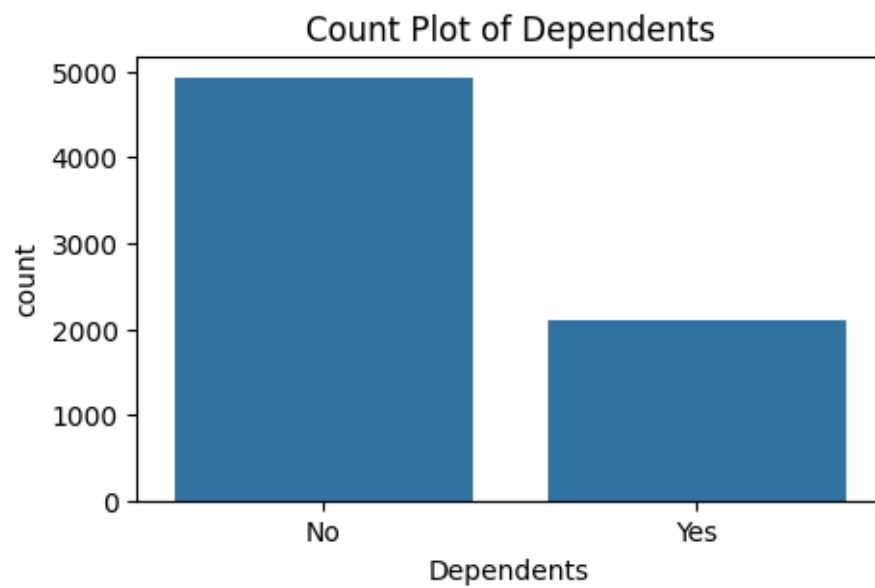

7 Count plot categorical columns

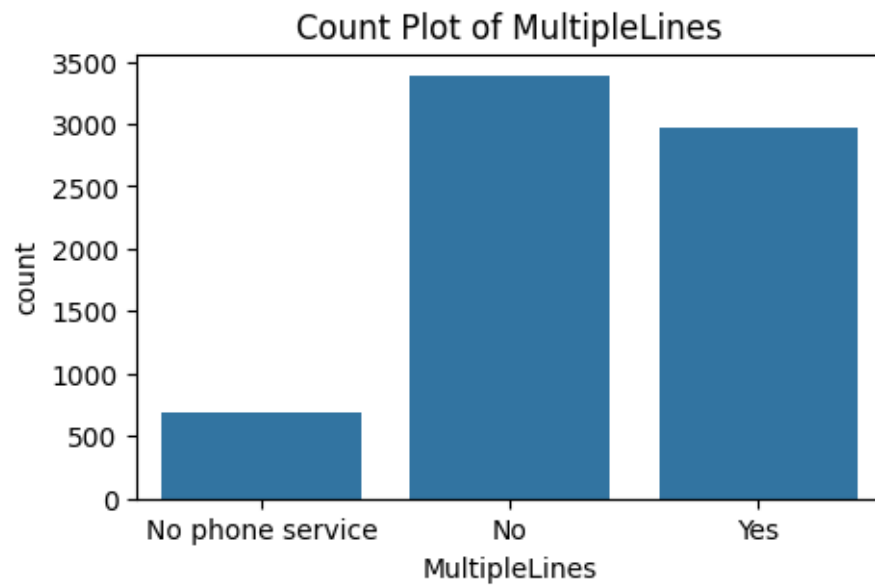

```
[39]: object_cols = df.select_dtypes(include="object").columns.to_list()



object_cols = ["SeniorCitizen"] + object_cols

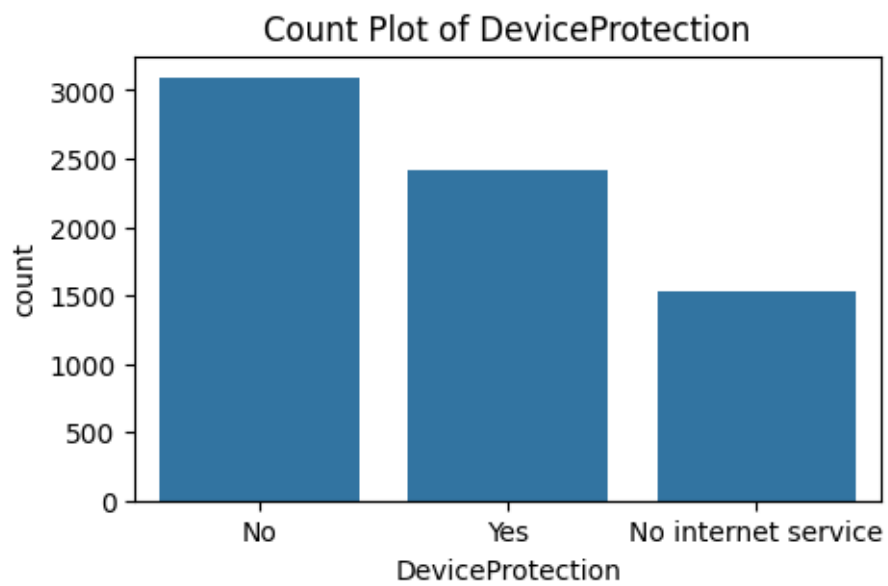

for col in object_cols:
    plt.figure(figsize=(5, 3))
    sns.countplot(x=df[col])
    plt.title(f"Count Plot of {col}")
    plt.show()
```






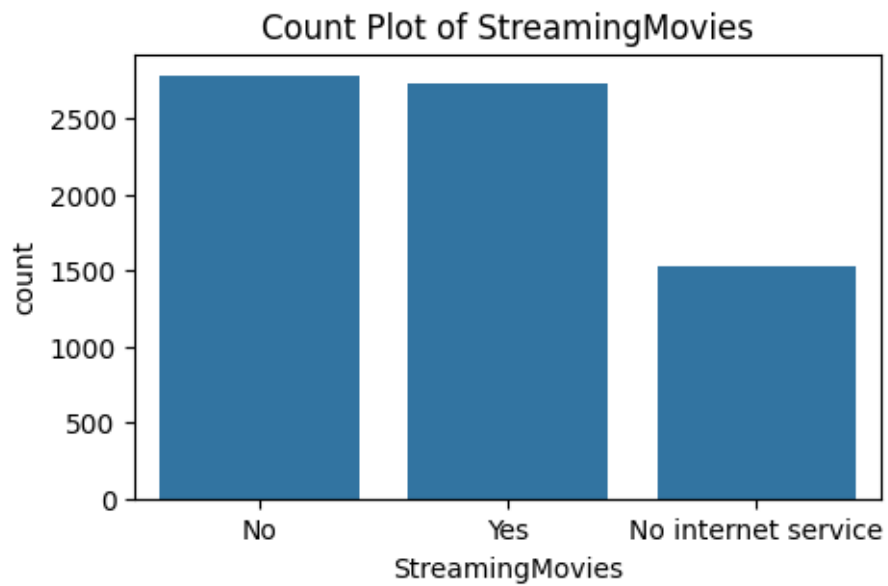



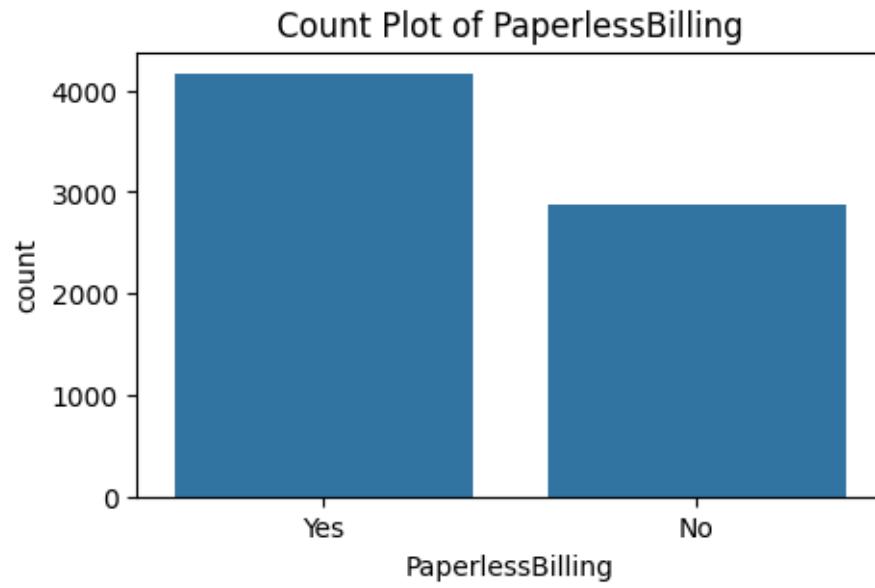















8 Data Processing

```
[40]: df.head(3)
```

```
[40]:   gender  SeniorCitizen  Partner  Dependents  tenure  PhoneService  \
0  Female              0     Yes         No         1             No
1   Male              0     No          No        34             Yes
2   Male              0     No          No         2             Yes

      MultipleLines  InternetService  OnlineSecurity  OnlineBackup  \
0  No phone service              DSL              No             Yes
1              No              DSL              Yes             No
2              No              DSL              Yes             Yes

      DeviceProtection  TechSupport  StreamingTV  StreamingMovies  Contract  \
0              No              No              No              No  Month-to-month
1              Yes              No              No              No      One year
2              No              No              No              No  Month-to-month

      PaperlessBilling  PaymentMethod  MonthlyCharges  TotalCharges  Churn
0              Yes  Electronic check         29.85         29.85     No
1              No    Mailed check         56.95        1889.50     No
2              Yes    Mailed check         53.85         108.15     Yes
```

9 Label encoding of target column

```
[41]: df["Churn"] = df["Churn"].replace({"Yes": 1, "No": 0})
```

```
C:\Users\Admin\AppData\Local\Temp\ipykernel_30384\2364848822.py:1:
FutureWarning: Downcasting behavior in `replace` is deprecated and will be
removed in a future version. To retain the old behavior, explicitly call
`result.infer_objects(copy=False)`. To opt-in to the future behavior, set
`pd.set_option('future.no_silent_downcasting', True)`
df["Churn"] = df["Churn"].replace({"Yes": 1, "No": 0})
```

```
[42]: df.head(4)
```

```
[42]:
```

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	\
0	Female	0	Yes	No	1	No	
1	Male	0	No	No	34	Yes	
2	Male	0	No	No	2	Yes	
3	Male	0	No	No	45	No	

	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	\
0	No phone service		DSL	No	Yes
1	No		DSL	Yes	No
2	No		DSL	Yes	Yes
3	No phone service		DSL	Yes	No

	DeviceProtection	TechSupport	StreamingTV	StreamingMovies	Contract	\
0	No	No	No	No	Month-to-month	
1	Yes	No	No	No	One year	
2	No	No	No	No	Month-to-month	
3	Yes	Yes	No	No	One year	

	PaperlessBilling	PaymentMethod	MonthlyCharges	TotalCharges	\
0	Yes	Electronic check	29.85	29.85	
1	No	Mailed check	56.95	1889.50	
2	Yes	Mailed check	53.85	108.15	
3	No	Bank transfer (automatic)	42.30	1840.75	

	Churn
0	0
1	0
2	1
3	0

```
[43]: print(df["Churn"].value_counts())
```

```
Churn
0    5174
1    1869
```

Name: count, dtype: int64

10 Label encoding

```
[79]: # identifying columns with object data type
object_columns = df.select_dtypes(include="object").columns
```

```
[80]: print(object_columns)
```

```
Index(['gender', 'Partner', 'Dependents', 'PhoneService', 'MultipleLines',
       'InternetService', 'OnlineSecurity', 'OnlineBackup', 'DeviceProtection',
       'TechSupport', 'StreamingTV', 'StreamingMovies', 'Contract',
       'PaperlessBilling', 'PaymentMethod'],
      dtype='object')
```

```
[81]: # initialize a dictionary to save the encoders
encoders = {}

# apply label encoding and store the encoders
for column in object_columns:
    label_encoder = LabelEncoder()
    df[column] = label_encoder.fit_transform(df[column])
    encoders[column] = label_encoder

# save the encoders to a pickle file
with open("encoders.pkl", "wb") as f:
    pickle.dump(encoders, f)
```

```
[82]: encoders
```

```
[82]: {'gender': LabelEncoder(),
       'Partner': LabelEncoder(),
       'Dependents': LabelEncoder(),
       'PhoneService': LabelEncoder(),
       'MultipleLines': LabelEncoder(),
       'InternetService': LabelEncoder(),
       'OnlineSecurity': LabelEncoder(),
       'OnlineBackup': LabelEncoder(),
       'DeviceProtection': LabelEncoder(),
       'TechSupport': LabelEncoder(),
       'StreamingTV': LabelEncoder(),
       'StreamingMovies': LabelEncoder(),
       'Contract': LabelEncoder(),
       'PaperlessBilling': LabelEncoder(),
       'PaymentMethod': LabelEncoder()}
```



```
[84]: df.head()
```

```
[84]:   gender  SeniorCitizen  Partner  Dependents  tenure  PhoneService  \
0        0             0        1           0         1             0
1        1             0        0           0        34             1
2        1             0        0           0         2             1
3        1             0        0           0        45             0
4        0             0        0           0         2             1

   MultipleLines  InternetService  OnlineSecurity  OnlineBackup  \
0              1                0                0              2
1              0                0                2              0
2              0                0                2              2
3              1                0                2              0
4              0                1                0              0

   DeviceProtection  TechSupport  StreamingTV  StreamingMovies  Contract  \
0                  0            0            0                0          0
1                  2            0            0                0          1
2                  0            0            0                0          0
3                  2            2            0                0          1
4                  0            0            0                0          0

   PaperlessBilling  PaymentMethod  MonthlyCharges  TotalCharges  Churn
0                  1              2           29.85         29.85     0
1                  0              3           56.95        1889.50     0
2                  1              3           53.85         108.15     1
3                  0              0           42.30        1840.75     0
4                  1              2           70.70         151.65     1
```

11 Training and Test data split

```
[44]: # splitting the features and target
X = df.drop(columns=["Churn"])
y = df["Churn"]
```

```
[45]: # split training and test data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=42)
```

```
[46]: print(y_train.shape)

(5634,)
```

```
[47]: print(y_train.value_counts())
```

```
Churn
0    4138
```

```
1    1496
Name: count, dtype: int64
```

12 (SMOTS)

```
[58]: smote = SMOTE(random_state=42)
```

```
[60]: from sklearn.preprocessing import LabelEncoder

# Encode categorical columns
categorical_cols = X.select_dtypes(include='object').columns
label_encoders = {}

for col in categorical_cols:
    le = LabelEncoder()
    X[col] = le.fit_transform(X[col])
    label_encoders[col] = le # save encoders in case you need to decode later

# Re-split the data after encoding
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪random_state=42)

# Apply SMOTE
from imblearn.over_sampling import SMOTE
smote = SMOTE(random_state=42)
X_train_smote, y_train_smote = smote.fit_resample(X_train, y_train)
```

```
[61]: X_train_smote, y_train_smote = smote.fit_resample(X_train, y_train)
```

```
[62]: print(y_train_smote.shape)
```

```
(8276,)
```

```
[63]: print(y_train_smote.value_counts())
```

```
Churn
0    4138
1    4138
Name: count, dtype: int64
```

13 Data Modelling

```
[64]: # dictionary of models
models = {
    "Decision Tree": DecisionTreeClassifier(random_state=42),
    "Random Forest": RandomForestClassifier(random_state=42),
    "XGBoost": XGBClassifier(random_state=42)
```

```
}
```

```
[65]: # dictionary to store the cross validation results
cv_scores = {}

# perform 5-fold cross validation for each model
for model_name, model in models.items():
    print(f"Training {model_name} with default parameters")
    scores = cross_val_score(model, X_train_smote, y_train_smote, cv=5,
    ↪scoring="accuracy")
    cv_scores[model_name] = scores
    print(f"{model_name} cross-validation accuracy: {np.mean(scores):.2f}")
    print("-"*70)
```

Training Decision Tree with default parameters

Decision Tree cross-validation accuracy: 0.78

Training Random Forest with default parameters

Random Forest cross-validation accuracy: 0.84

Training XGBoost with default parameters

XGBoost cross-validation accuracy: 0.83

```
[66]: cv_scores
```

```
[66]: {'Decision Tree': array([0.68297101, 0.71601208, 0.81993958, 0.83564955,
0.83746224]),
      'Random Forest': array([0.72826087, 0.7734139 , 0.90332326, 0.89969789,
0.8978852 ]),
      'XGBoost': array([0.71135266, 0.74864048, 0.91178248, 0.88640483, 0.91117825])}
```

14 Random Forest give highest accuracy

```
[67]: rfc = RandomForestClassifier(random_state=42)
```

```
[68]: rfc.fit(X_train_smote, y_train_smote)
```

```
[68]: RandomForestClassifier(random_state=42)
```

```
[69]: print(y_test.value_counts())
```

Churn

0 1036

1 373

Name: churn, dtype: int64

15 Model Evaluation

```
[71]: # evaluate on test data
y_test_pred = rfc.predict(X_test)

print("Accuracy Score:\n", accuracy_score(y_test, y_test_pred))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_test_pred))
print("Classification Report:\n", classification_report(y_test, y_test_pred))
```

Accuracy Score:

0.7771469127040455

Confusion Matrix:

[[879 157]

[157 216]]

Classification Report:

	precision	recall	f1-score	support
0	0.85	0.85	0.85	1036
1	0.58	0.58	0.58	373
accuracy			0.78	1409
macro avg	0.71	0.71	0.71	1409
weighted avg	0.78	0.78	0.78	1409

```
[72]: # save the trained model as a pickle file
model_data = {"model": rfc, "features_names": X.columns.tolist()}

with open("customer_churn_model.pkl", "wb") as f:
    pickle.dump(model_data, f)
```

16 Load the saved model and build a predictive system

```
[73]: # load the saved model and the feature names

with open("customer_churn_model.pkl", "rb") as f:
    model_data = pickle.load(f)

loaded_model = model_data["model"]
feature_names = model_data["features_names"]
```

```
[74]: print(loaded_model)
```

RandomForestClassifier(random_state=42)

```
[75]: print(feature_names)
```

```
['gender', 'SeniorCitizen', 'Partner', 'Dependents', 'tenure', 'PhoneService',
'MultipleLines', 'InternetService', 'OnlineSecurity', 'OnlineBackup',
'DeviceProtection', 'TechSupport', 'StreamingTV', 'StreamingMovies', 'Contract',
'PaperlessBilling', 'PaymentMethod', 'MonthlyCharges', 'TotalCharges']
```

```
[85]: input_data = {
    'gender': 'Female',
    'SeniorCitizen': 0,
    'Partner': 'Yes',
    'Dependents': 'No',
    'tenure': 1,
    'PhoneService': 'No',
    'MultipleLines': 'No phone service',
    'InternetService': 'DSL',
    'OnlineSecurity': 'No',
    'OnlineBackup': 'Yes',
    'DeviceProtection': 'No',
    'TechSupport': 'No',
    'StreamingTV': 'No',
    'StreamingMovies': 'No',
    'Contract': 'Month-to-month',
    'PaperlessBilling': 'Yes',
    'PaymentMethod': 'Electronic check',
    'MonthlyCharges': 29.85,
    'TotalCharges': 29.85
}

input_data_df = pd.DataFrame([input_data])

with open("encoders.pkl", "rb") as f:
    encoders = pickle.load(f)

# encode categorical features using the saved encoders
for column, encoder in encoders.items():
    input_data_df[column] = encoder.transform(input_data_df[column])

# make a prediction
prediction = loaded_model.predict(input_data_df)
pred_prob = loaded_model.predict_proba(input_data_df)

print(prediction)

# results
print(f"Prediction: {'Churn' if prediction[0] == 1 else 'No Churn'}")
print(f"Prediction Probability: {pred_prob}")
```

```
[0]
Prediction: No Churn
Prediciton Probability: [[0.79 0.21]]
```

```
[86]: encoders
```

```
[86]: {'gender': LabelEncoder(),
      'Partner': LabelEncoder(),
      'Dependents': LabelEncoder(),
      'PhoneService': LabelEncoder(),
      'MultipleLines': LabelEncoder(),
      'InternetService': LabelEncoder(),
      'OnlineSecurity': LabelEncoder(),
      'OnlineBackup': LabelEncoder(),
      'DeviceProtection': LabelEncoder(),
      'TechSupport': LabelEncoder(),
      'StreamingTV': LabelEncoder(),
      'StreamingMovies': LabelEncoder(),
      'Contract': LabelEncoder(),
      'PaperlessBilling': LabelEncoder(),
      'PaymentMethod': LabelEncoder()}
```

17 To do:

1. Implement Hyperparameter Tuining
2. Try Model Selection
3. Try downsampling
4. Try to address teh overfitting
5. Try Startified k fold CV

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```